

# Parallel Processing in Python

Presented by:  
Ong Chin Hwee (@ongchinhwee)

27 August 2019  
Python User Group Singapore, Aug 2019 Meetup

# About me

---

- Current role: Data Engineer at ST Engineering
- Background in aerospace engineering and computational modelling
- Experience working on aerospace-related projects in collaboration with academia and industry partners
- Find me if you would like to chat about Industry 4.0 and flight + travel!

# A typical data science workflow

---

1. Define problem objective
2. Data collection and pipeline
3. Data parsing/preprocessing and Exploratory Data Analysis (EDA)
4. Feature engineering
5. Model training
6. Model evaluation
7. Visualization and Reporting
8. Model deployment

What do you think are some of the bottlenecks in a data science project?

---

# Bottlenecks in a data science project

---

- Lack of data / Poor quality data
- Data Preprocessing
  - The 80/20 data science dilemma
    - In reality, it's closer to 90/10
- The organization itself

# Bottlenecks in a data science project

---

- Data Preprocessing
  - Pandas faces low performance and long runtime issues when dealing with large datasets ( $> 1$  GB)

# Bottlenecks in a data science project

---

- Data Preprocessing
  - Pandas faces low performance and long runtime issues when dealing with large datasets ( $> 1$  GB)
  - Slow loops in Python
    - Loops are run on the interpreter, not compiled (unlike loops in C)

# Scope of Talk

---

I will talk about:

1. What is parallel processing?



# Scope of Talk

---

I will talk about:

1. What is parallel processing?
2. Synchronous vs asynchronous execution

# Scope of Talk

---

I will talk about:

1. What is parallel processing?
2. Synchronous vs asynchronous execution
3. When should you go for parallelism?

# Scope of Talk

---

I will talk about:


1. What is parallel processing?
2. Synchronous vs asynchronous execution
3. When should you go for parallelism?
4. Parallel processing in Python

# Scope of Talk

---

I will NOT talk about:

1. Parallel computing hardware



# Scope of Talk

---

I will **NOT** talk about:

1. Parallel computing hardware
2. Apache Spark, MapReduce etc.

The word "spark" is written in a black, cursive, handwritten style. To the right of the word is a small, orange, multi-pointed starburst or asterisk symbol.

# Scope of Talk

---

I will NOT talk about:

1. Parallel computing hardware
2. Apache Spark, MapReduce etc.
3. The GIL controversy

Did some pythonic developer just say



**THREADS?**

# What is parallel processing?

Let's imagine I own a bakery cafe.

---



# Task 1: Toast 100 slices of bread

Assumptions:

1. I'm using single-slice toasters.  
(Yes, they actually exist.)
2. Each slice of toast takes 2 minutes to make.
3. No overhead time.




Image taken from:

<https://www.mitsubishielectric.co.jp/home/breadoven/product/to-st1-t/feature/index.html>

# Sequential Processing



---



= 25 bread slices

# Sequential Processing


---




= 25 bread slices


Processor/Worker:  
Toaster

# Sequential Processing




 = 25 bread slices

**Processor/Worker:**  
Toaster

 = 25 toasts

# Sequential Processing


---



**Execution Time** = 100 toasts × 2 minutes/toast  
= **200 minutes**

# Parallel Processing


---



 = 25 bread slices


# Parallel Processing

---



# Parallel Processing

---




Processor (Core):  
Toaster



# Parallel Processing

---




**Processor (Core):**  
Toaster

Task is executed using  
a **pool** of **4 toaster**  
**subprocesses**.

Each toasting  
subprocess runs **in**  
**parallel** and  
**independently** from  
each other.


# Parallel Processing



**Processor (Core):**  
Toaster

Output of each  
toasting process is  
**consolidated** and  
**returned** as an overall  
output (which may or  
may not be ordered).

# Parallel Processing



**Execution Time**  
= 100 toasts  $\times$  2  
minutes/toast  $\div$   
4 toasters  
= **50 minutes**

**Speedup**  
= **4 times**

# Synchronous vs Asynchronous Execution

---

# Task 2: Brew 100 cups of coffee

Assumptions:


1. I am doing BOTH tasks 1 and 2 (toast + coffee).
2. One coffee maker to make one cup of coffee.
3. Each cup of coffee takes 5 minutes to make.



Image taken from: <https://www.crateandbarrel.com/breville-barista-espresso-machine/s267619>

# Synchronous Execution


---



Process 1: Toast a slice of  
bread on single-slice toaster  
Duration: 2 minutes

# Synchronous Execution

---




Process 1: Toast a slice of bread on single-slice toaster  
Duration: 2 minutes

Process 2: Brew a cup of coffee on coffee machine  
Duration: 5 minutes



# Synchronous Execution



Process 1: Toast a slice of bread on single-slice toaster  
Duration: 2 minutes

Process 2: Brew a cup of coffee on coffee machine  
Duration: 5 minutes



Output: **1 toast + 1 coffee**

**Total Execution Time = 2 minutes + 5 minutes = 7 minutes**




# Asynchronous Execution

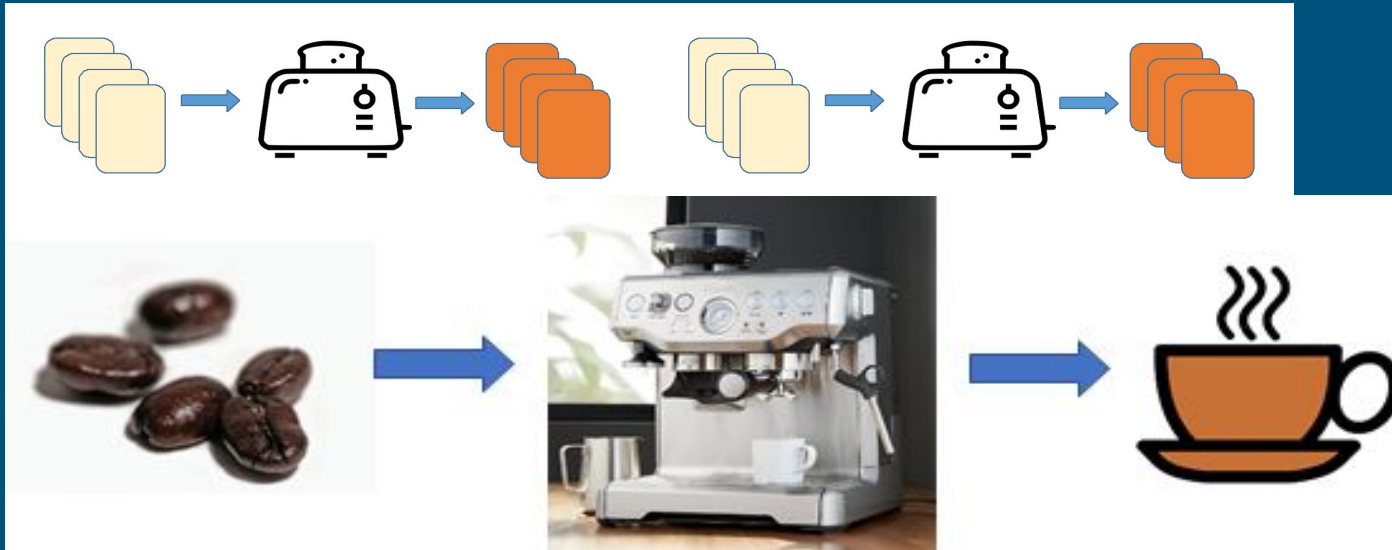
While brewing coffee:



Make some toasts:



# Asynchronous Execution



Output: 2 toasts + 1 coffee (1 more toast!)  
Total Execution Time = 5 minutes

When is it a good idea to go for  
parallelism?

---

# When is it a good idea to go for parallelism?

---

(or, “Is it a good idea to simply buy a 256-core processor and parallelize all your codes?”)

# Practical Considerations

---

- Is your code already optimized?
  - Sometimes, you might need to rethink your approach.
  - Example: Use list comprehensions instead of for-loops for array iterations.

# Practical Considerations

---

- Is your code already optimized?
  - Sometimes, you might need to rethink your approach.
- Problem architecture
  - Nature of problem limits how successful parallelization can be.
  - If your problem consists of processes which depend on each others' outputs, maybe not.

# Practical Considerations

---

- Is your code already optimized?
  - Sometimes, you might need to rethink your approach.
- Problem architecture
  - Nature of problem limits how successful parallelization can be.
- Overhead in parallelism
  - There will always be parts of the work that cannot be parallelized. → **Amdahl's Law**
  - Extra time required for coding and debugging (parallelism vs sequential code)

# Amdahl's Law and Parallelism

---

**Amdahl's Law** states that the theoretical speedup is defined by the fraction of code  $p$  that can be parallelized:

$$S = \frac{1}{(1 - p) + \frac{p}{N}}$$

$S$ : Theoretical speedup (theoretical latency)

$p$ : Fraction of the code that can be parallelized


$N$ : Number of processors (cores)



# Amdahl's Law and Parallelism

If there are no parallel parts ( $p = 0$ ): Speedup = 0


If all parts are parallel ( $p = 1$ ):



# Amdahl's Law and Parallelism

If there are no parallel parts ( $p = 0$ ): Speedup = 0

If all parts are parallel ( $p = 1$ ):  
Speedup  $\rightarrow \infty$




# Amdahl's Law and Parallelism

If there are no parallel parts ( $p = 0$ ): Speedup = 0

If all parts are parallel ( $p = 1$ ):  
Speedup  $\rightarrow \infty$

Speedup is limited by **fraction of the work that is not parallelizable** - will not improve even with infinite number of processors



# Parallel Processing in Python Programming

---

# Why is parallelism so tricky in Python?

---

## Global Interpreter Lock (GIL)

In CPython, the GIL is a **mutex** (mutually exclusive object) that allows **only one thread** to hold control of the Python interpreter **at once**.

Necessary mainly because CPython's memory management is **not fully thread-safe**.

While non-parallel codes can run faster with the GIL, it does not allow parallel thread execution.

# Why is parallelism so tricky in Python?

---

## Global Interpreter Lock (GIL)

However, stuff that happens outside the GIL realm (I/O and libraries written on external C codes) can **bypass the GIL**.

# How to do Parallel Processing in Python?

---

# Parallel Processing in Python

---

## `concurrent.futures` module

- High-level API for launching asynchronous parallel tasks
- Introduced in Python 3.2 as an abstraction layer over **multiprocessing** module
- Two modes of execution:
  - *ThreadPoolExecutor()* for multithreading
  - *ProcessPoolExecutor()* for multiprocessing




# Multiprocessing vs Multithreading

---

## Multiprocessing:

System allows executing multiple processes at the same time using multiple processors



# Multiprocessing vs Multithreading

---

## **Multiprocessing:**

System allows executing multiple processes at the same time using multiple processors

## **Multithreading:**

System executes multiple threads of sub-processes at the same time within a single processor

# Multiprocessing vs Multithreading

---

## Multiprocessing:

System allows executing multiple processes at the same time using multiple processors

Better option for **processing large volumes of data**

## Multithreading:

System executes multiple threads of sub-processes at the same time within a single processor

Best suited for **I/O operations**

# ProcessPoolExecutor vs ThreadPoolExecutor

---

From the Python Standard Library documentation:

For *ProcessPoolExecutor*, this method chops iterables into a number of chunks which it submits to the pool as **separate tasks**. The (approximate) size of these chunks can be specified by **setting chunksize to a positive integer**. For very long iterables, using a large value for chunksize can significantly improve performance compared to the default size of 1. With *ThreadPoolExecutor*, chunksize has **no effect**.

# Recap: map() in Python Built-in Functions

---

**map()** takes as input:

1. The function that you would like to run, and
2. A list (iterable) where each element of the list is a single input to that function;

and returns an iterator that **yields** the results of the function being applied to every element of the list.

# map() in concurrent.futures

Similarly, **executor.map()** takes as input:

1. The function that you would like to run, and
2. A list (iterable) where each element of the list is a single input to that function;

and returns an iterator that **yields** the results of the function being applied to every element of the list.

# Practical Implementation

---

# Case: Image Processing

---

Dataset: Shopee National Data Science Challenge  
(<https://www.kaggle.com/c/ndsc-advanced>)

Size: **77.6GB** of image files

Data Quality: Images in the dataset are of **different formats** (some are RGB while others are RGBA) and **different dimensions**



# Image Resizing Code

---

```
from PIL import Image
```

```
def image_proc(index):  
    '''Convert + resize image'''  
    im = Image.open(define_imagepath(index))  
    im = im.convert("RGB")  
    im_resized = np.array(im.resize((64,64)))  
  
    return im_resized
```

# Batch Processing Code

---

```
def arraypartition_calc(start, batch_size):  
    '''Process images in partition/batch'''  
    end = start + batch_size  
    if end > N:  
        end = N  
  
    partition_list = [image_proc(image) for image  
in range(start, end)]  
  
    return partition_list
```

# Without Parallelism

---

```
import sys
import time

N = 35000 # size of dataset to be processed
start = 0
batch_size = 1000
partition = int(np.ceil(N/step))
partition_count = 0
imagearray_list = [None] * partition

start_cpu_time = time.clock()
start_wall_time = time.time()
```

# Without Parallelism

---

```
while start < N:

    end = start + batch_size
    if end > N:
        end = N

    imagearray_list[partition_count] =
[arraypartition_calc(image) for image in range(start, end)]

    start += batch_size
    partition_count += 1
```

# Without Parallelism

---

```
while start < N:

    end = start + batch_size
    if end > N:
        end = N

    imagearray_list[partition_count] =
[arraypartition_calc(image) for image in range(start, end)]

    start += batch_size
    partition_count += 1
```

# Without Parallelism

---

```
while start < N:

    end = start + batch_size
    if end > N:
        end = N

    imagearray_list[partition_count] =
[arraypartition_calc(image) for image in range(start, end)]

    start += batch_size
    partition_count += 1
```

## Execution Speed:

3300 images after 7 hours  
= 471.43 images/hr

# With Parallelism

---

```
N = 35000
start = 0
batch_size = 1000
partition, mod = divmod(N, batch_size)

if mod:
    partition_index = [i * batch_size for i in range(start //
batch_size, partition + 1)]
else:
    partition_index = [i * batch_size for i in range(start //
batch_size, partition)]
```

# With Parallelism

---

```
import sys
import time
from concurrent.futures import ProcessPoolExecutor

start_cpu_time = time.clock()
start_wall_time = time.time()

with ProcessPoolExecutor() as executor:
    future = executor.map(arraypartition_calc, partition_index)

imgarray_np = np.array([x for x in future])
```



# With Parallelism

---

## Execution Speed:

35000 images after 3.6 hours  
= 9722.22 images/hr

```
import sys
import time
from concurrent.futures import ProcessPoolExecutor

start_cpu_time = time.clock()
start_wall_time = time.time()

with ProcessPoolExecutor() as executor:
    future = executor.map(arraypartition_calc, partition_index)

imgarray_np = np.array([x for x in future])
```

# With Parallelism

---

```
import sys
import time
from concurrent.futures import ProcessPoolExecutor
```

```
start_cpu_time = time.clock()
start_wall_time = time.time()
```

```
with ProcessPoolExecutor() as executor:
    future = executor.map(arraypartition_calc, partition_index)

imgarray_np = np.array([x for x in future])
```

**Execution Speed:**

35000 images after 3.6 hours  
= 9722.22 images/hr

**Speedup: 19.4 times**

# With Parallelism

---

```
import sys
import time
from concurrent.futures import ProcessPoolExecutor

start_cpu_time = time.clock()
start_wall_time = time.time()

with ProcessPoolExecutor() as executor:
    future = executor.map(arraypartition_calc, partition_index)

    imgarray_np = np.array([x for x in future])
```

Extract results from  
iterator (similar to  
generator)

# Key Takeaways

---

# Parallel Processing in Python

---

**concurrent.futures** module:

- Provides **hassle-free, high-level** implementation of parallel processing without delving into parallelization architectures
- Part of **Python Standard Library** since Python 3.2 - no need for additional installations!

# Parallel Processing in Python

---

- Not all processes should be parallelized
  - Amdahl's Law on parallelism
  - Extra time required for coding and debugging (parallelism vs sequential code)
  - If the cost of rewriting your code for parallelization outweighs the time savings from parallelizing your code (especially if your process only takes a few hours), maybe you should consider **other ways of optimizing your code** instead.

# References

---

Official Python documentation on concurrent.futures  
(<https://docs.python.org/3/library/concurrent.futures.html>)

Built-in Functions - Python 3.7.4 Documentation  
(<https://docs.python.org/3/library/functions.html#map>)

First Programs and How to Think in CUDA. (2011). CUDA Application Design and Development, 1–31.  
doi:10.1016/b978-0-12-388426-8.00001-x

# Contact

---

**Ong Chin Hwee**

LinkedIn: ongchinhwee

Twitter: @ongchinhwee

<https://ongchinhwee.me>





# Women Who Code Connect Asia 2019



WOMEN WHO  
**CODE**  
CONNECT ASIA

CONNECTING WOMEN  
WITH PROFESSIONAL  
OPPORTUNITIES TO  
ADVANCE THEIR CAREERS

AUGUST 31ST

SINGAPORE

#CONNECTASIA

# Women Who Code Connect Asia 2019

---

Date: August 31st 2019 (Saturday)

Time: 8:00 AM – 6:00 PM SGT

Venue: Lifelong Learning Institute

For anyone purchasing using Singapore Dollars:

Member price - \$75 USD : \$100 SGD

Open Collective link:

<https://opencollective.com/wwcodeconnectasia/events/women-who-code-connect-asia-2019-august-31-st-member-tickets-50595ev>